

Problem Solving by Search

1

CHAPTER 3

Stuart Russell and Peter Norvig, Artificial Intelligence: A
Modern Approach, Global Edition 3/E

Outline

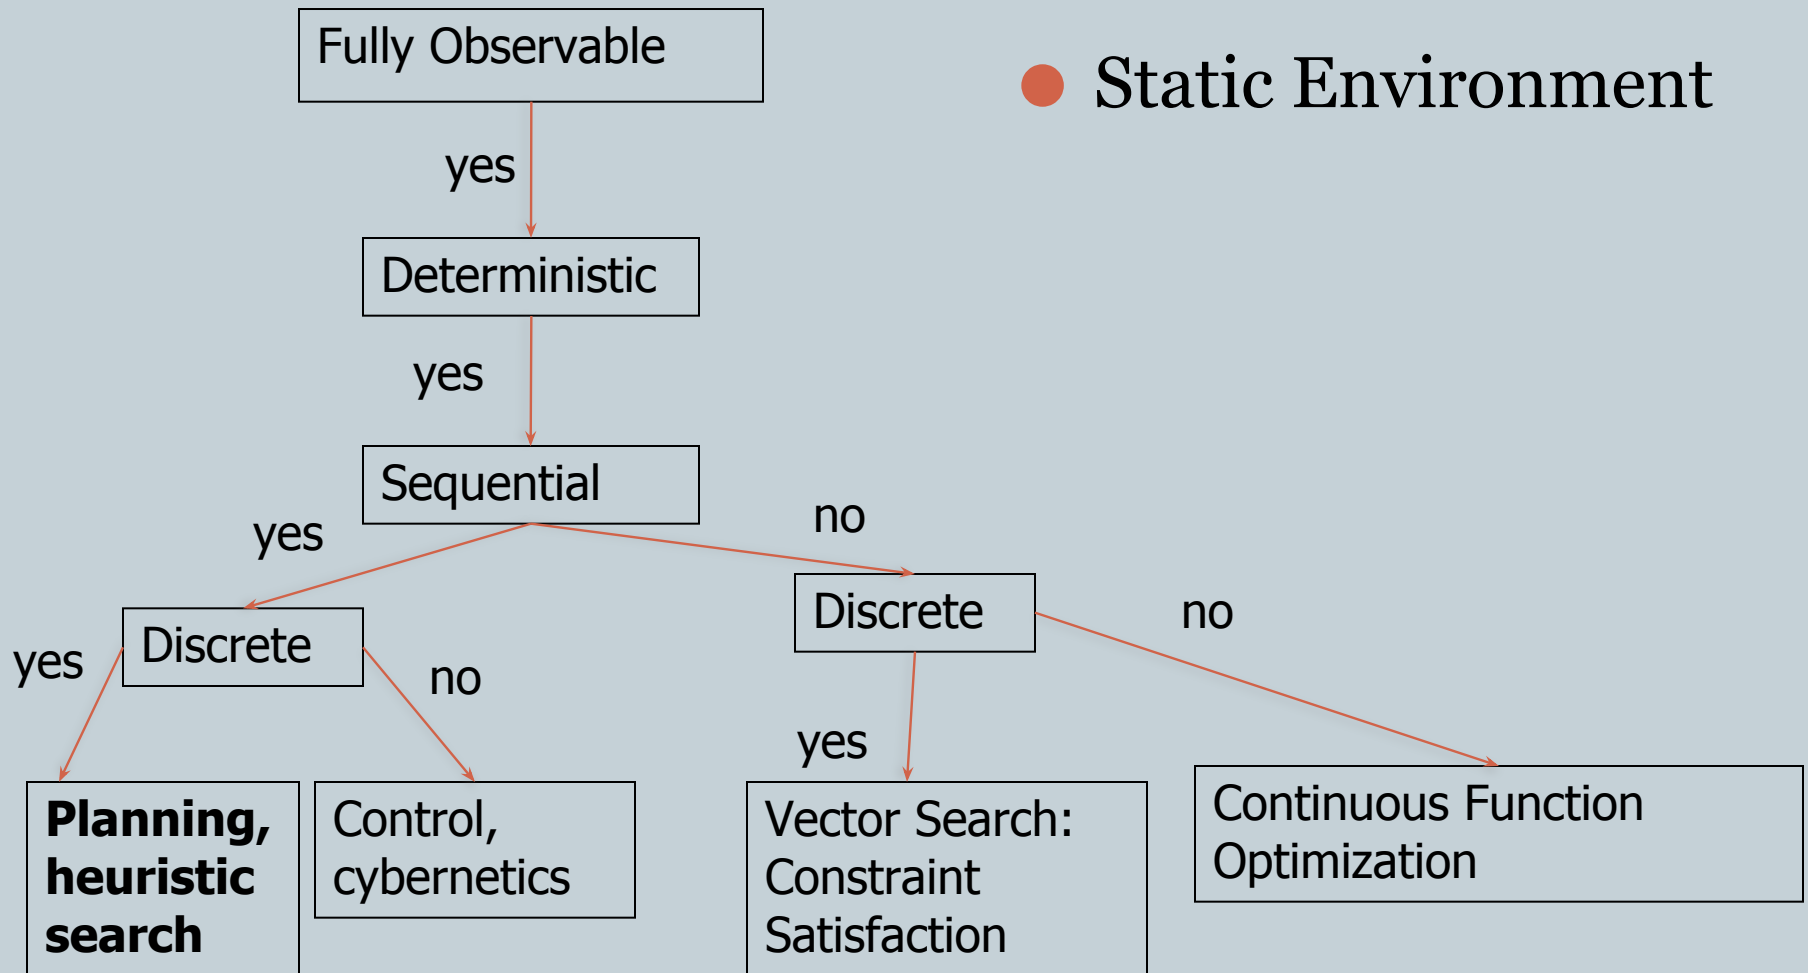
2

- Problem formulation: representing sequential problems.
- Example problems.
- Planning for solving sequential problems without uncertainty.
- Basic search algorithms

Environment Type Discussed In this Lecture

3

● Static Environment



Choice in a Deterministic Known Environment

4

- Without uncertainty, choice is trivial in principle: choose what you know to be the best option.
- Trivial if the problem is represented in a look-up table.

Option	Value
Chocolate	10
Coffee	20
Book	15

This is the standard problem representation in decision theory (economics).

Computational Choice Under Certainty

5

- But choice can be *computationally* hard if the problem information is represented differently.
- Options may be **structured** and the best option needs to be constructed.
 - E.g., an option may consist of a path, sequence of actions, plan, or strategy.
- The value of options may be given **implicitly** rather than explicitly.
 - E.g., cost of paths need to be computed from map.

Problem Types


6

- **Deterministic**, fully observable -> single-state problem
 - Agent knows exactly which state it will be in; solution is a **sequence**
- **Non-observable** -> conformant problem
 - Agent may have no idea where it is; solution (if any) is a **sequence**
- **Nondeterministic** and/or partially observable -> contingency problem
 - percepts provide new information about current state
solution is a contingent plan or a policy
often **interleave search, execution**
- **Unknown state space** -> exploration problem (“online”)

Sequential Action Example

7

- **Deterministic, fully observable: single-state problem**
 - Agent knows exactly which state it will be in; solution is a sequence
 - Vacuum world: everything observed
 - Romania: The full map is observed
- **Single-state:**
Start in #5. Solution??
 - [Right, Suck]



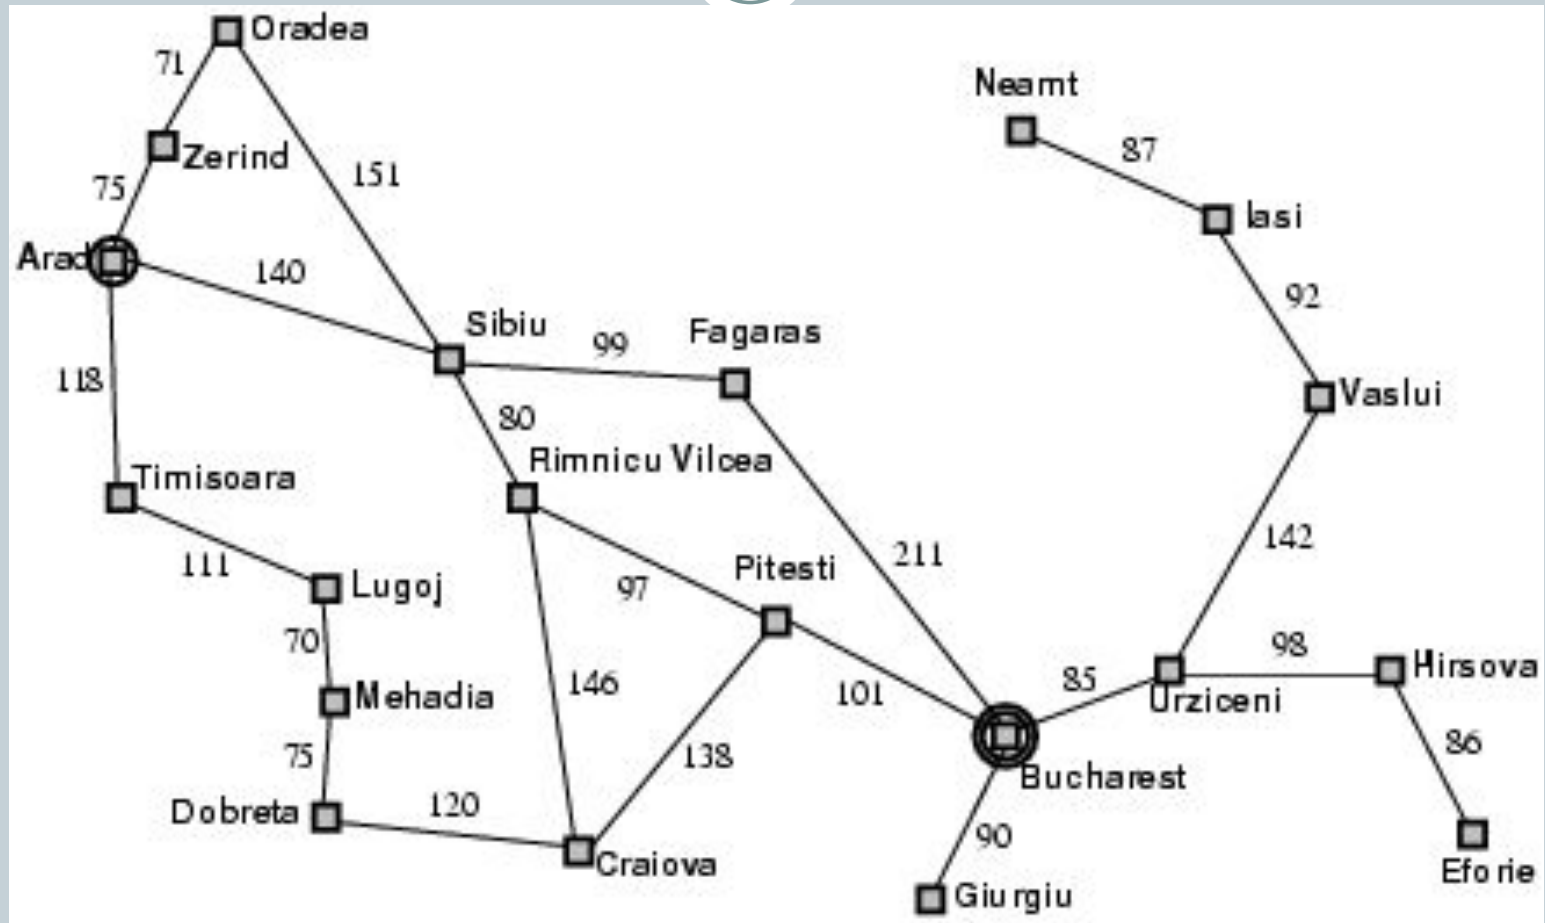
Example: Romania

8

- On holiday in Romania; currently in Arad.
- Formulate goal:
 - be in Bucharest
- Formulate problem:
 - **states**: various cities
 - **actions**: drive between cities
- Find solution:
 - sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

Example: Romania

9



Abstraction: The process of removing details from a representation
Is the map a good representation of the problem? What is a good replacement?

Single-state problem formulation

10

A problem is defined by 4 items:

- **initial state** e.g., “at Arad”
- **Successor function** $S(x)$ = set of action–state
- **Goal test**, can be
 - explicit, e.g., x = “at Bucharest”, or “checkmate” in chess
 - implicit, e.g., $\text{NoDirt}(x)$
- **Path cost** (additive) e.g., sum of distances, number of actions executed, etc. $c(x, a, y)$ is the step cost, assumed to be ≥ 0

A solution is a sequence of actions leading from the initial state to a goal state

The successor function

11

- Successor function: for a given state, returns a set of action/new-state pairs.
- Vacuum-cleaner world: $(A, \text{dirty}, \text{clean}) \rightarrow ('Left', (A, \text{dirty}, \text{clean})), ('Right', (B, \text{dirty}, \text{clean})), ('Suck', (A, \text{clean}, \text{dirty})), ('NoOp', (A, \text{dirty}, \text{clean}))$
- Romania: $\text{In}(\text{Arad}) \rightarrow ((\text{Go}(\text{Timisoara}), \text{In}(\text{Timisoara}), (\text{Go}(\text{Sibiu}), \text{In}(\text{Sibiu}))), (\text{Go}(\text{Zerind}), \text{In}(\text{Zerind})))$

Size of space

12

- 8-puzzle: $9!/2 = 181,000$ states (easy)
- 15-puzzle: ~ 1.3 trillion states (pretty easy)
- 24-puzzle: $\sim 10^{25}$ states (hard)
- TSP, 20 cities: $20! = 2.43 \times 10^{18}$ states (hard)


Selecting a state space

13

- Real world is complex
 - state space must be **abstracted** for problem solving
- (Abstract) state = set of real states
- (Abstract) action = complex combination of real actions
 - e.g., "Arad -> Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- (Abstract) solution =
 - set of real paths that are solutions in the real world
- Each abstract action should be "easier" than the original problem

Vacuum world state space graph


14



- states?
- actions?
- goal test?
- path cost?

Vacuum world state space graph

15



- states? integer dirt and robot location
- actions? *Left, Right, Suck*
- goal test? no dirt at all locations
- path cost? 1 per action

Example: The 8-puzzle

16

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- states?
- actions?
- goal test?
- path cost?

Example: The 8-puzzle

17

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8


Goal State

- states? locations of tiles
- actions? move blank left, right, up, down
- goal test? = goal state (given)
- path cost? 1 per move

[Note: optimal solution of n -Puzzle family is NP-hard]

Example: robotic assembly

18



- states?: real-valued coordinates of robot joint angles parts of the object to be assembled
- actions?: continuous motions of robot joints
- goal test?: complete assembly
- path cost?: time to execute

Problem-solving agents

19

```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  static: seq, an action sequence, initially empty
          state, some description of the current world state
          goal, a goal, initially null
          problem, a problem formulation

  state ← UPDATE-STATE(state, percept)
  if seq is empty then
    goal ← FORMULATE-GOAL(state)
    problem ← FORMULATE-PROBLEM(state, goal)
    seq ← SEARCH(problem)
  action ← FIRST(seq)
  seq ← REST(seq)
  return action
```

Note: this is offline problem solving; solution executed “eyes closed.”

Tree search algorithms

20

- Basic idea:

- offline, simulated exploration of state space by generating successors of already-explored states (a.k.a. ~expanding states)

function TREE-SEARCH(*problem*, *strategy*) **returns** a solution, or failure

 initialize the search tree using the initial state of *problem*

loop do

 if there are no candidates for expansion **then return** failure

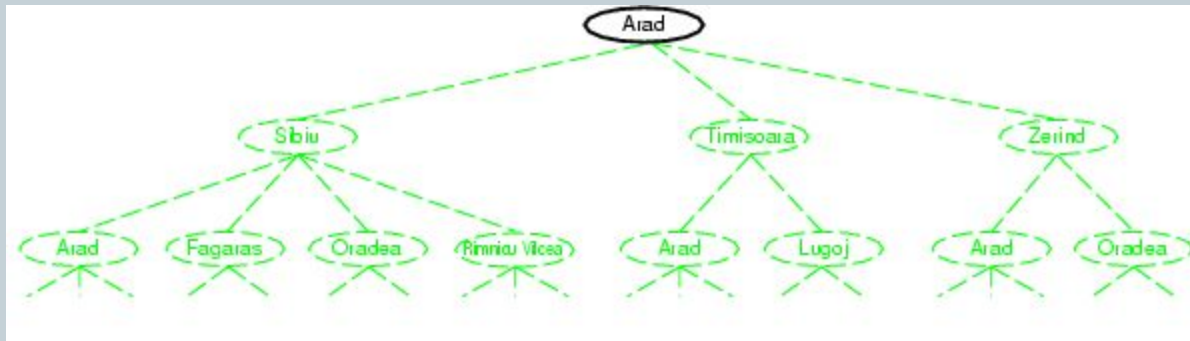
 choose a leaf node for expansion according to *strategy*

 if the node contains a goal state **then return** the corresponding solution

 else expand the node and add the resulting nodes to the search tree

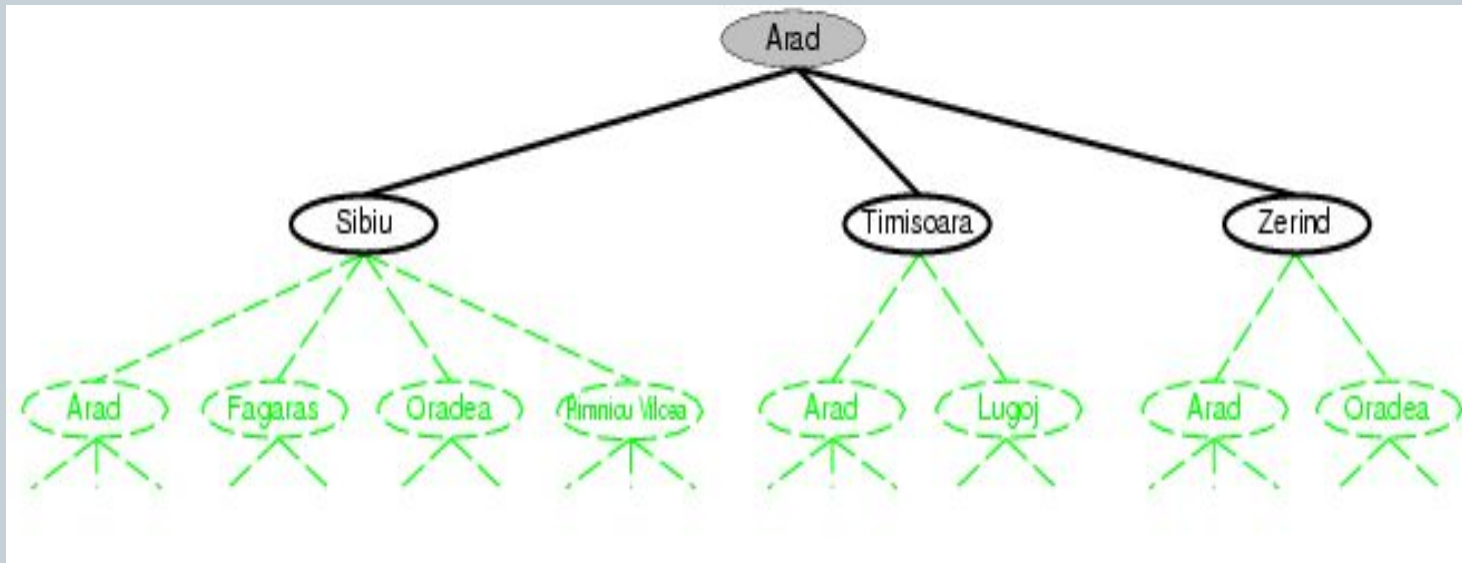
Tree search example

21



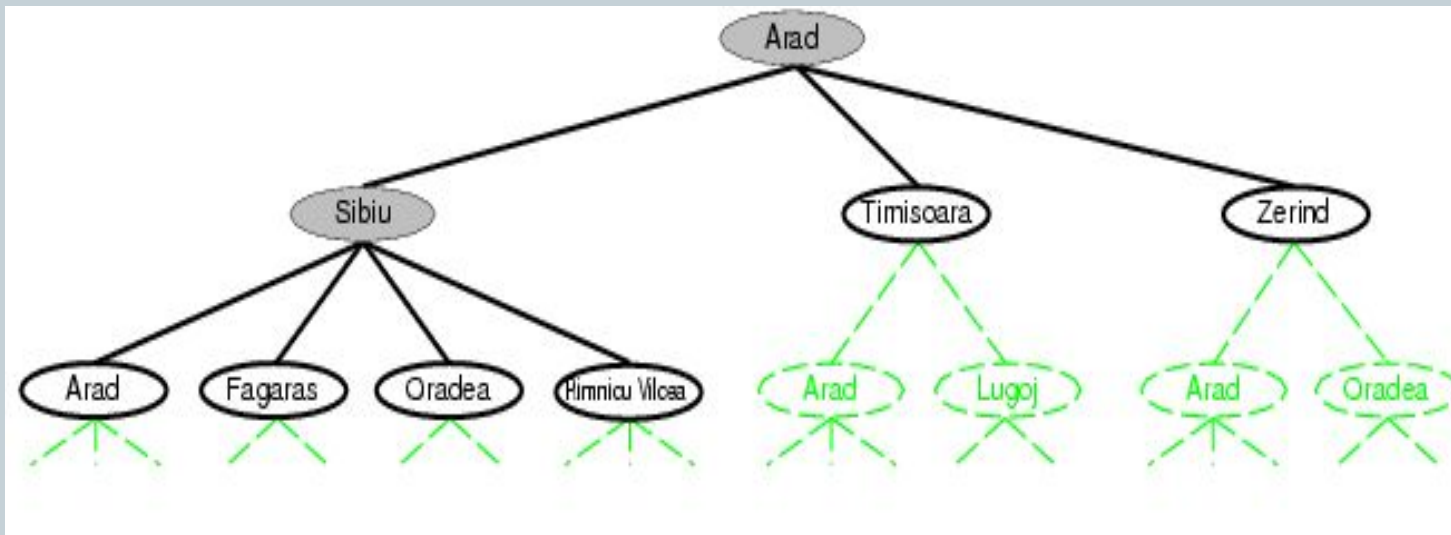
Tree search example

22



Tree search example

23



Search Graph vs. State Graph

24

- Be careful to distinguish
 - Search tree: nodes are **sequences of actions**.
 - State Graph: Nodes are states of the environment.
 - We will also consider soon **search graphs**.
- Demo: <http://aispace.org/search/>

Search strategies

25

- A search strategy is defined by picking the **order of node expansion**
- Strategies are evaluated along the following dimensions:
 - **completeness**: does it always find a solution if one exists?
 - **time complexity**: number of nodes generated
 - **space complexity**: maximum number of nodes in memory
 - **optimality**: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
 - b : maximum branching factor of the search tree
 - d : depth of the least-cost solution
 - m : maximum depth of the state space (may be ∞)

Search Strategies

26

- Uninformed (blind) search
- Informed Search
- Adversarial Search (Game Theory)

Uninformed search strategies


27

- **Uninformed** search strategies use only the information available in the problem definition
- Uninformed search (blind search)
 - Breadth-first search
 - Depth-first search
 - Depth-limited search
 - Iterative deepening search

Breadth-first search

28


- Expand shallowest unexpanded node
- Implementation:
 - is a FIFO queue, i.e., new successors go at end



Breadth-first search

29


- Expand shallowest unexpanded node
- **Implementation:**
 - is a FIFO queue, i.e., new successors go at end



Breadth-first search

30


- Expand shallowest unexpanded node
- **Implementation:**
 - is a FIFO queue, i.e., new successors go at end



Breadth-first search

31

- Expand shallowest unexpanded node
<http://aispace.org/search/>
- **Implementation:**
 - is a FIFO queue, i.e., new successors go at end



Properties of breadth-first search


32

- Complete? Time? Space? Optimal?
- Complete? Yes (if b is finite)
- Time? $1+b+b^2+b^3+\dots +b^d + b(b^d-1) = O(b^{d+1})$
- Space? $O(b^{d+1})$ (keeps every node in memory)
- Optimal? Yes (if cost = 1 per step)
- **Space** is the bigger problem (more than time)

Depth-first search

33


- Expand deepest unexpanded node
- **Implementation:**
 - LIFO queue, i.e., put successors at front



Depth-first search

34


- Expand deepest unexpanded node
- **Implementation:**
 - LIFO queue, i.e., put successors at front



Depth-first search

35


- Expand deepest unexpanded node
- **Implementation:**
 - LIFO queue, i.e., put successors at front



Depth-first search

36


- Expand deepest unexpanded node
- **Implementation:**
 - *frontier* = LIFO queue, i.e., put successors at front



Depth-first search

37


- Expand deepest unexpanded node
- **Implementation:**
 - LIFO queue, i.e., put successors at front



Depth-first search

38


- Expand deepest unexpanded node
- **Implementation:**
 - LIFO queue, i.e., put successors at front



Depth-first search

39


- Expand deepest unexpanded node
- **Implementation:**
 - LIFO queue, i.e., put successors at front



Depth-first search

40


- Expand deepest unexpanded node
- **Implementation:**
 - LIFO queue, i.e., put successors at front



Depth-first search

41


- Expand deepest unexpanded node
- **Implementation:**
 - LIFO queue, i.e., put successors at front



Depth-first search

42

- Expand deepest unexpanded node
- **Implementation:**
 - LIFO queue, i.e., put successors at front




43

-

Depth-first search

44

- Expand deepest unexpanded node
<http://aispace.org/search/>
- **Implementation:**
 - LIFO queue, i.e., put successors at front



Properties of depth-first search

45

- Complete? Time? Space? Optimal?
- Complete? No: fails in infinite-depth spaces, spaces with loops
 - Modify to avoid repeated states along path (graph search)
 - complete in finite spaces
- Time? $O(b^m)$: terrible if maximum depth m is much larger than solution depth d
 - but if solutions are dense, may be much faster than breadth-first
- Space? $O(bm)$, i.e., linear space! Store single path with unexpanded siblings.
 - Seems to be common in animals and humans.
- Optimal? No.
- Important for exploration (on-line search).

Depth-limited search

46

- depth-first search with depth limit l ,
 - i.e., nodes at depth l have no successors
 - Solves infinite loop problem
- Common AI strategy: let user choose search/resource bound.
Complete? No if $l < d$:
- Time? $O(b^l)$:
- Space? $O(bl)$, i.e., linear space!
- Optimal? No if $l > b$

Iterative deepening search

47

function ITERATIVE-DEEPENING-SEARCH(*problem*) **returns** a solution, or failure

inputs: *problem*, a problem

for *depth* \leftarrow 0 **to** ∞ **do**

result \leftarrow DEPTH-LIMITED-SEARCH(*problem*, *depth*)

if *result* \neq cutoff **then return** *result*

Iterative deepening search $l = 0$

48





Limit = 0



Iterative deepening search $l = 1$

49









Limit = 1



Iterative deepening search $l = 2$

50

Limit = 2



51

Iterative deepening search

52

- Number of nodes generated in a depth-limited search to depth d with branching factor b :

$$N_{DLS} = b^0 + b^1 + b^2 + \dots + b^{d-2} + b^{d-1} + b^d$$

- Number of nodes generated in an iterative deepening search to depth d with branching factor b :

$$N_{IDS} = (d+1)b^0 + d b^1 + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$

- For $b = 10$, $d = 5$,

$$N_{DLS} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$$

$$N_{IDS} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$$

- Overhead = $(123,456 - 111,111)/111,111 = 11\%$

Properties of iterative deepening search

53

- Complete? Yes
- Time? $(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- Space? $O(bd)$
- Optimal? Yes, if step cost = 1

Summary of algorithms

54

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Time	b^d	b^d	b^m	b^l	b^d	$b^{d/2}$
Space	b^d	b^d	bm	bl	bd	$b^{d/2}$
Optimal?	Yes	Yes	No	No	Yes	Yes
Complete?	Yes	Yes	No	Yes, if $l \geq d$	Yes	Yes

Graph search


55

```
function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    if STATE[node] is not in closed then
      add STATE[node] to closed
      fringe ← INSERTALL(EXPAND(node, problem), fringe)
```


- Simple solution: just keep track of which states you have visited.
- Usually easy to implement in modern computers.

The Separation Property of Graph Search


56



(a)



(b)



(c)

- Black: expanded nodes.
- White: frontier nodes.
- Grey: unexplored nodes.

Summary

57

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored
- Variety of uninformed search strategies
- Iterative deepening search uses only linear space and not much more time than other uninformed algorithms



End of Chapter 3

Informed search algorithms



CHAPTER 4

Stuart Russell and Peter Norvig, Artificial Intelligence: A
Modern Approach, Global Edition 3/E

Outline

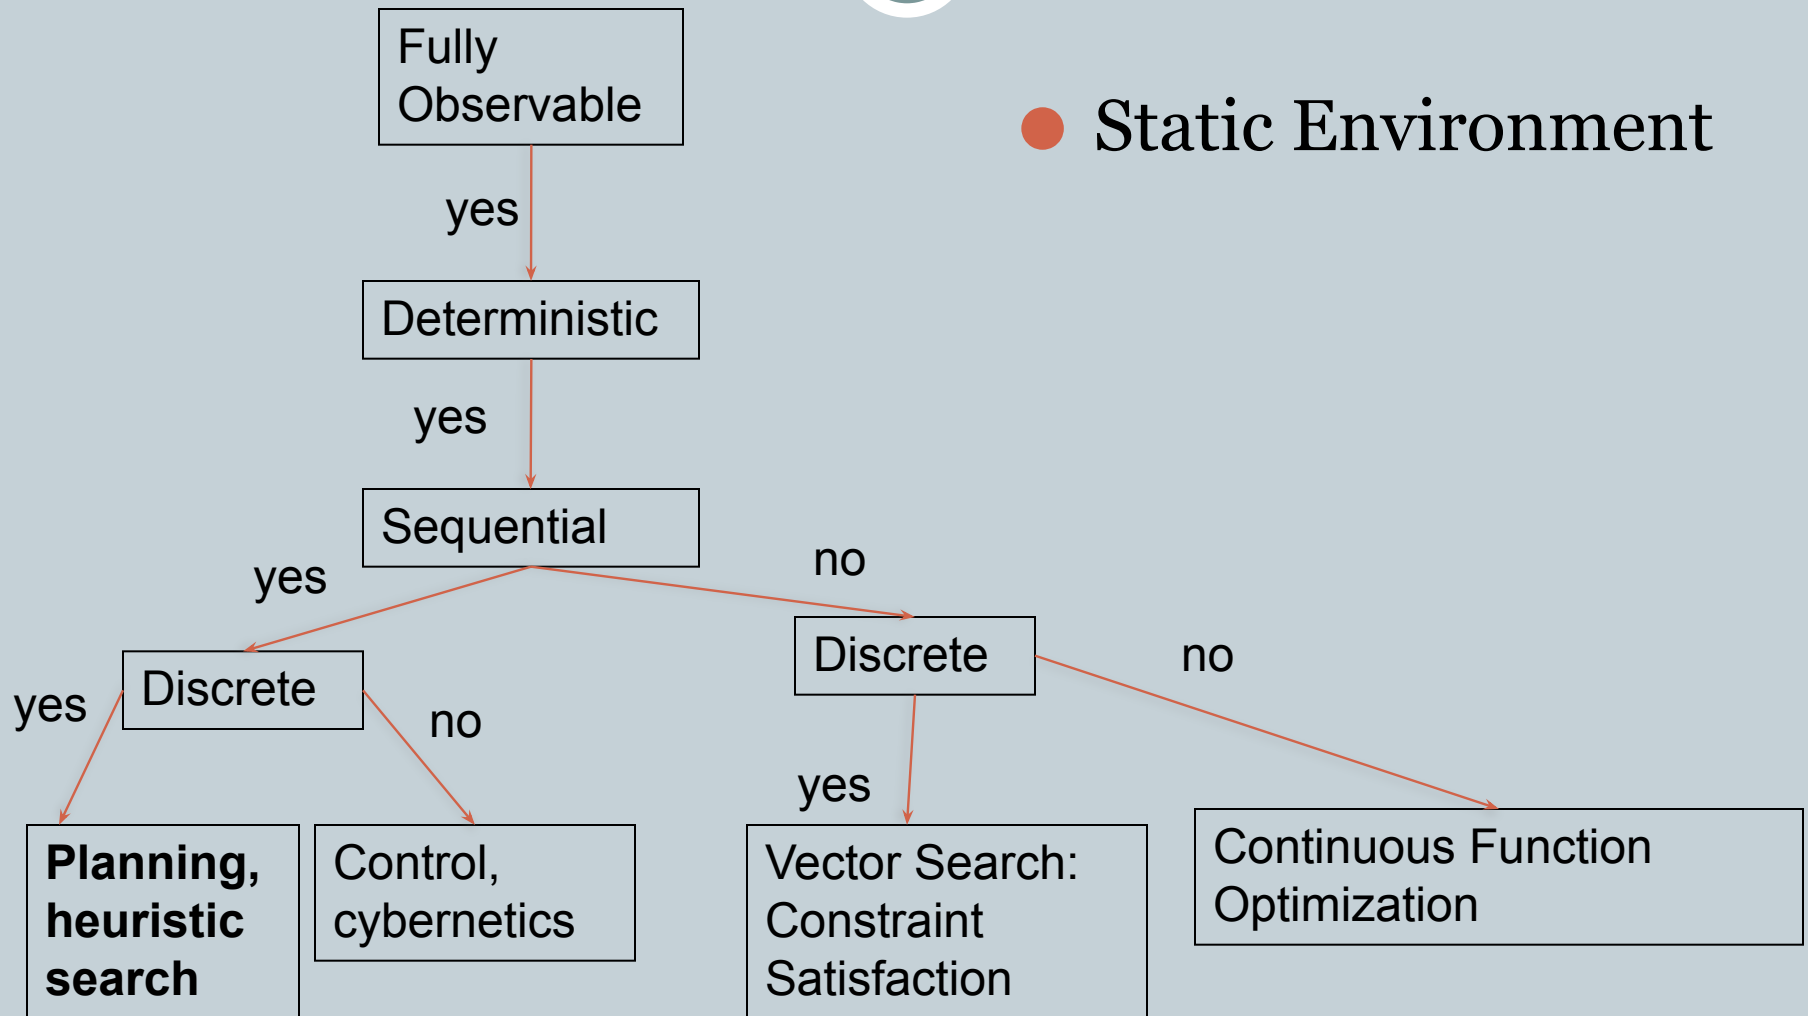
60

- Best-first search
- A^* search
- Heuristics
- Local search algorithms
- Hill-climbing search
- Simulated annealing search
- Local beam search

Environment Type Discussed In this Lecture

61

● Static Environment



Review: Tree search

62

```
function TREE-SEARCH(problem, fringe) returns a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem] applied to STATE(node) succeeds return node
    fringe ← INSERTALL(EXPAND(node, problem), fringe)
```

- A search strategy is defined by picking the **order of node expansion**
- Which nodes to check first?

Knowledge and Heuristics

63

- Simon and Newell, *Human Problem Solving*, 1972.
- S&N: intelligence comes from **heuristics** that help find promising states fast.


Best-first search

64

- Idea: use an **evaluation function** $f(n)$ for each node
 - estimate of "desirability"
 - Expand most desirable unexpanded node
- Implementation:
 - Order the nodes in frontier in decreasing order of desirability
- Special cases:
 - greedy best-first search
 - A^* search

Romania with step costs in km

65



Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Greedy best-first search

66

- Evaluation function
 - $f(n) = h(n)$ (**h**euristic)
 - = estimate of cost from n to *goal*
- e.g., $h_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal

Greedy best-first search example

67

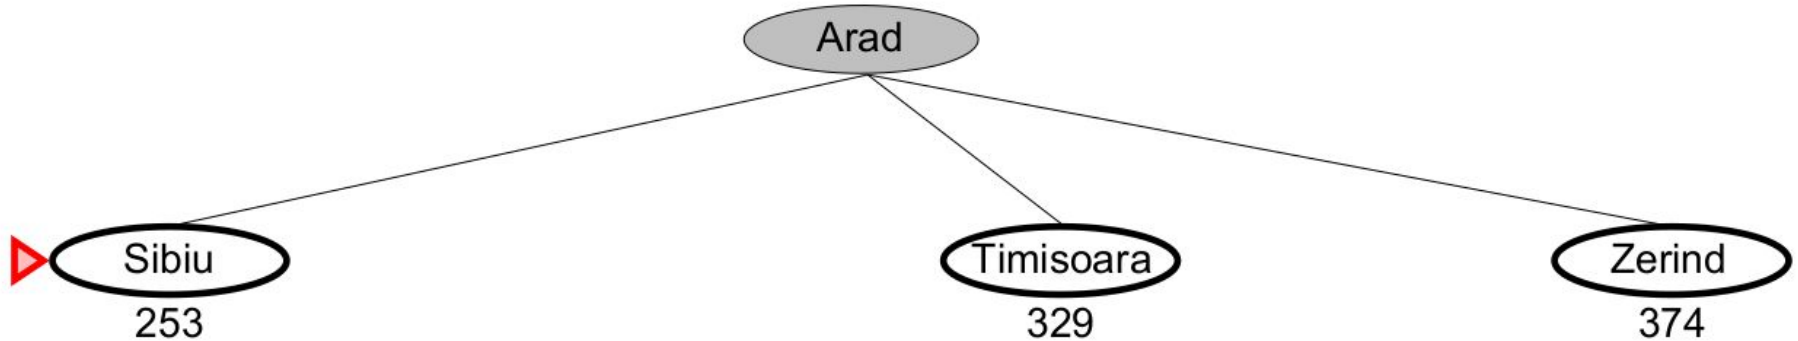


Arad

366


Greedy best-first search example

68




Greedy best-first search example

69



Greedy best-first search example

70



<http://aispace.org/search/>

Properties of greedy best-first search

71

- Complete? No – can get stuck in loops,
 - e.g. as Oradea as goal
 - Iasi □ Neamt □ Iasi □ Neamt □
- Time? $O(b^m)$, but a good heuristic can give dramatic improvement
- Space? $O(b^m)$ -- keeps all nodes in memory
- Optimal? No


A* search

72

- Idea: avoid expanding paths that are already expensive.
- Very important!
- Evaluation function $f(n) = g(n) + h(n)$
 $g(n)$ = cost so far to reach n
- $h(n)$ = estimated cost from n to goal
- **$f(n)$ = estimated total cost of path through n to goal**

A* search example

73

 Arad
 $366 = 0 + 366$

A* search example

74

Arad

Sibiu

$$393 = 140 + 253$$

Timisoara

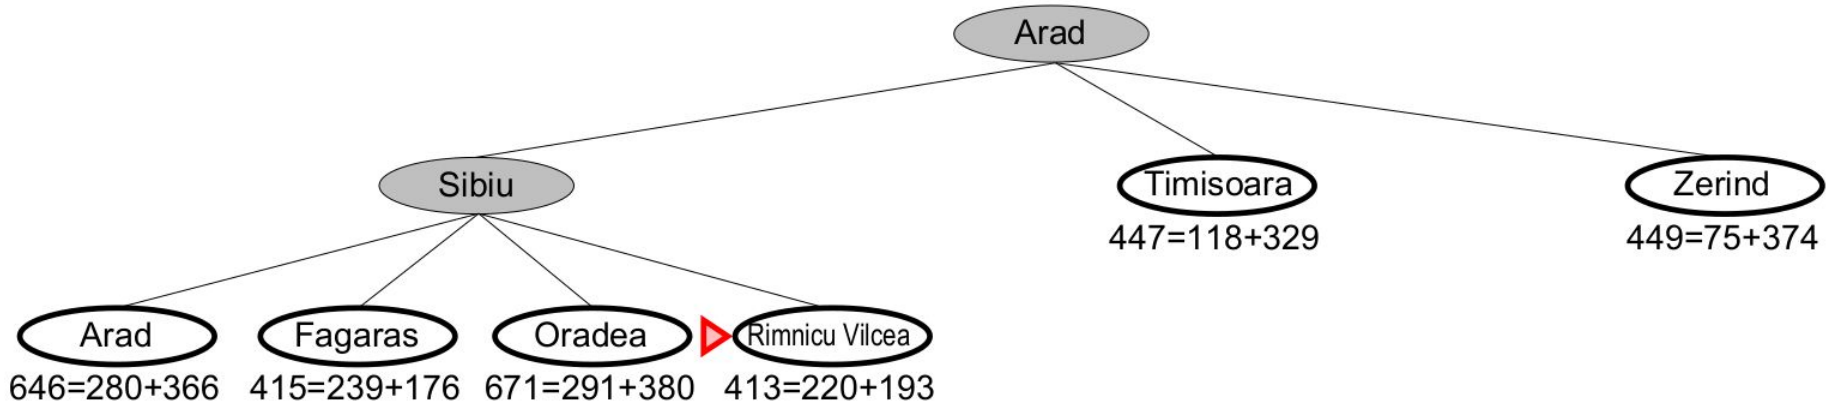
$$447 = 118 + 329$$

Zerind

$$449 = 75 + 374$$


A* search example

75



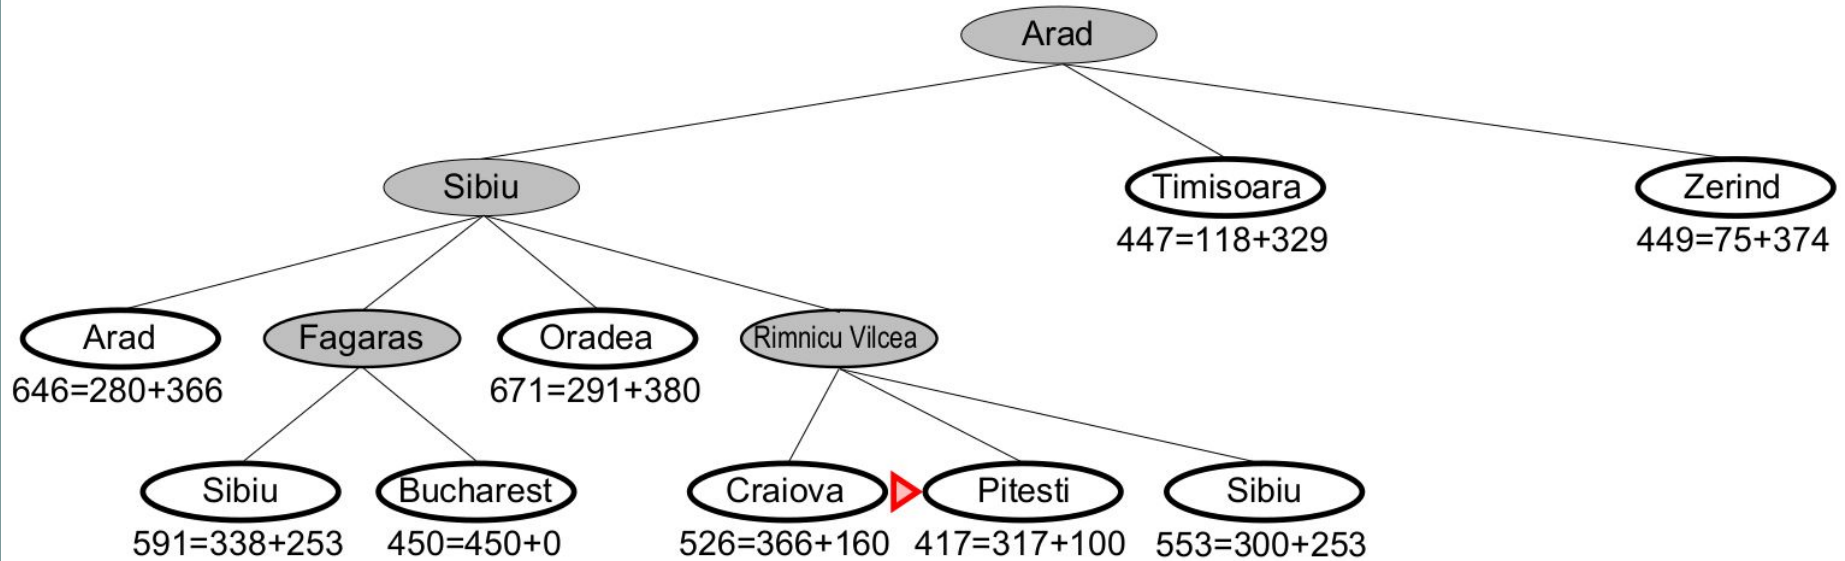
A* search example

76



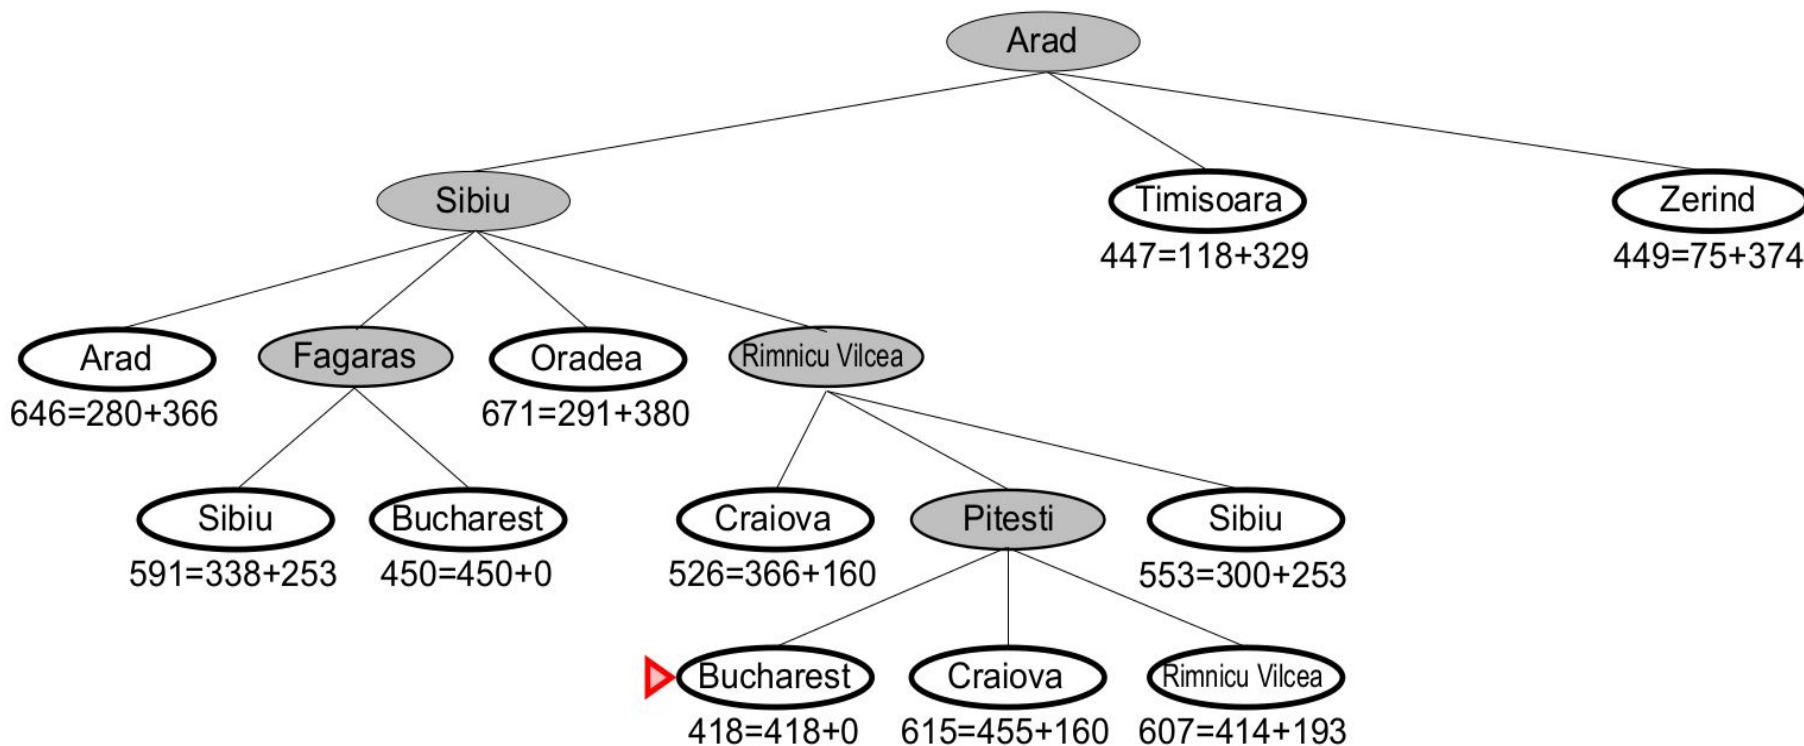
A* search example

77



A* search example

78



<http://aispace.org/search/>

- We stop when the node with the lowest f-value is a goal state.
- Is this guaranteed to find the shortest path?

Properties of A*

79

- Complete? Yes (unless there are infinitely many nodes with $f \leq f(G)$)
- Time? Exponential
- Space? Keeps all nodes in memory
- Optimal? Yes

Summary

80

- Heuristic functions estimate costs of shortest paths
- Good heuristics can dramatically reduce search cost
- Greedy best-first search expands lowest h
 - incomplete and not always optimal
- A* search expands lowest $g + h$
 - complete and optimal
 - also optimally efficient (up to tie-breaks)